

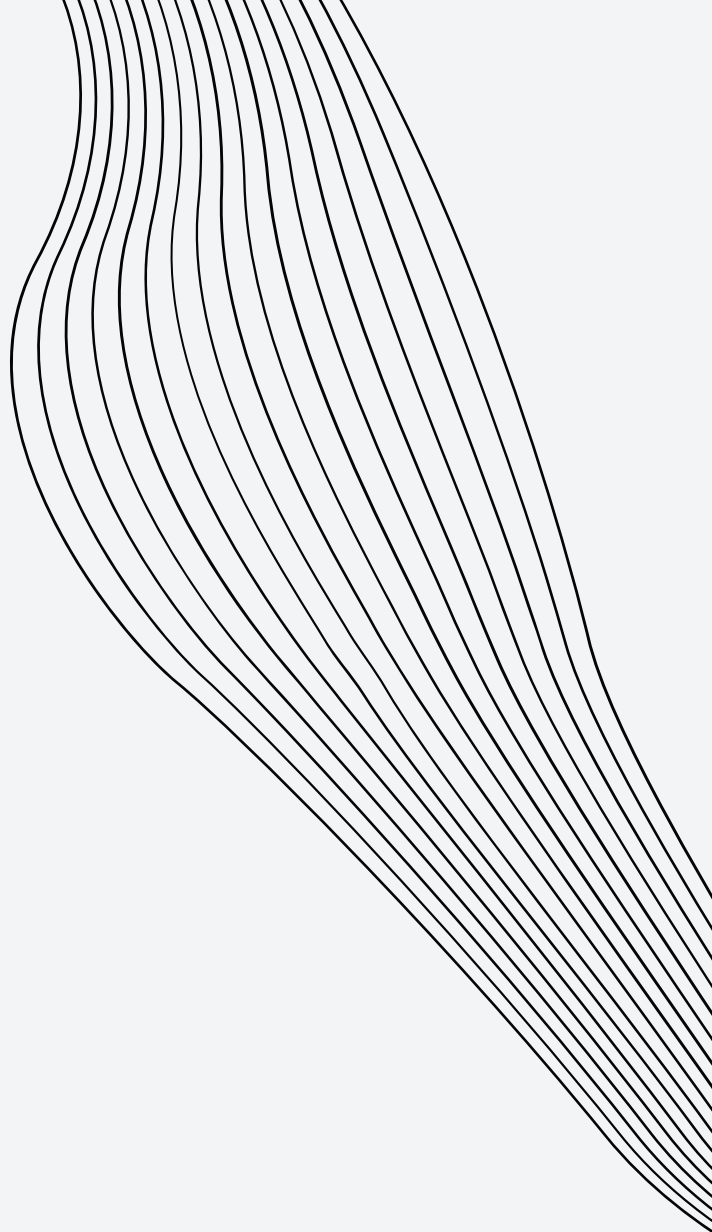
MULTILINGUAL DEPRESSION DETECTION USING SPEECH-BASED PRE-TRAINED MODEL

EERIK SVEN PUUDIST, FEDOR STOMAKHIN

RIO SCHULTS, ZUZANNA MAKOWSKA



EXPLAINABLE AUTOMATED MACHINE LEARNING - AUTUMN 2023

CONTENT

- 
- 
- 01** PROBLEM DESCRIPTION
 - 02** DATASETS
 - 03** FEATURE EXTRACTION
 - 04** AUTOML – HYPEROPT

PROJECT DESCRIPTION

**MULTILINGUAL DEPRESSION
DETECTION USING SPEECH-BASED
PRE-TRAINED MODEL**



DATASETS

ANDROID

reading

interview

Spanish

5 seconds

silence removed

labeled

EDAIC

test, train, validation

English

ANDROID

```
android_interview_depressed_count: 2208  
android_reading_depressed_count: 553  
android_interview_healthy_count: 2659  
android_reading_healthy_count: 479  
all samples: 5899
```

<https://drive.google.com/uc?id=1OQ87c6vEKkTuLu2Z3jYzOP6-pvCytojz>

EDAIC

```
edaic_train_count: 12547  
edaic_test_count: 2446  
edaic_validation_count: 2987  
all samples: 17980
```

<https://drive.google.com/uc?id=1PT9Iij7DJOB1s4iOT4gT3jZxpoZqSpzU>

EDAIC - LABELS

```
# train set
edaic_train_labels_dict = edaic_train_labels.set_index('ID')['Value'].to_dict()
edaic_train_recordings_names = os.listdir(edaic_train_path)
edaic_train_recordings_labels = []
for file in edaic_train_recordings_names:
    file_id = file.split('_')[0]
    if int(file_id) in edaic_train_labels_dict:
        edaic_train_recordings_labels.append((file, edaic_train_labels_dict[int(file_id)]))
    else:
        edaic_train_recordings_labels.append((file, None))
```



WHISPER


```
# Whisper- Base
from transformers import AutoFeatureExtractor, WhisperModel
# from datasets import load_dataset

model = WhisperModel.from_pretrained("openai/whisper-base")
feature_extractor = AutoFeatureExtractor.from_pretrained("openai/whisper-base")
model.to(device)
import torchaudio
def extract_features(path):
    sample_rate = 16000
    array, fs = torchaudio.load(path)
    input = feature_extractor(array.squeeze(), sampling_rate = sample_rate, return_tensors = 'pt')
    input = input.to(device)
    input = input.input_features
    with torch.no_grad():
        outputs = model.encoder(input)
    last_hidden_states = outputs.last_hidden_state.squeeze().mean(axis = 0).to(device).numpy()
    return last_hidden_states
```

```
android_reading_healthy_files = [os.path.join(android_reading_healthy_path, file) for file in os.listdir(android_reading_healthy_path)]

android_reading_healthy_features_whisper = []

for file in android_reading_healthy_files:
    features = extract_features(file)
    file_name = os.path.basename(file)
    android_reading_healthy_features_whisper.append([file_name, 0] + list(features))
```

```
def extract_features_with_labels(recordings_labels, base_path):
    features_labels_dataset = []
    for file, label in recordings_labels:
        file_path = os.path.join(base_path, file)
        features = extract_features(file_path)

        features_labels_dataset.append([file, label] + list(features))
    return features_labels_dataset

edaic_train_features_whisper = extract_features_with_labels(edaic_train_recordings_labels, edaic_train_path)

edaic_test_features_whisper = extract_features_with_labels(edaic_test_recordings_labels, edaic_test_path)

edaic_validation_whisper = extract_features_with_labels(edaic_validation_recordings_labels, edaic_validation_path)
```



BASELINE

ANDROID - DATA PREPARATION

```
android_reading_healthy_features_whisper_df = pd.DataFrame(android_reading_healthy_features_whisper)
android_reading_depressed_features_whisper_df = pd.DataFrame(android_reading_depressed_features_whisper)
android_interview_healthy_features_whisper_df = pd.DataFrame(android_interview_healthy_features_whisper)
android_interview_depressed_features_whisper_df = pd.DataFrame(android_interview_depressed_features_whisper)

android = android_reading_healthy_features_whisper_df.append([ android_reading_depressed_features_whisper_df, android_interview_healthy_features_whisper_df, android_interview_depressed_features_whisper_df])
android.reset_index(inplace = True)
android = android.drop('index', axis=1)
android
```

	0	1	2	3	4	5	6	7	8	9
0	38_CM27_3_4	0	-0.325004	-0.773003	0.336232	-0.826059	0.200828	1.663229	-0.456727	0.036573
1	44_CF37_3_5	0	-0.323559	-0.755329	0.209577	-0.853966	0.209243	1.502434	-0.385898	-0.245348
2	41_CM71_2_8	0	-0.316730	-0.754254	0.186741	-0.993816	0.174278	1.494282	-0.547302	0.042305
3	10_CF51_2_8	0	-0.251689	-0.507035	0.329449	-0.812264	0.245561	1.537125	-0.494353	-0.107126
4	38_CM27_3_2	0	-0.385444	-0.721874	0.129210	-1.014975	0.111982	1.407310	-0.593221	-0.001678

STRATEGY N°3

ANDROID - SPLITTING THE DATA

```
[ ] X = android.iloc[:,2:]  
    Y = android[1]
```

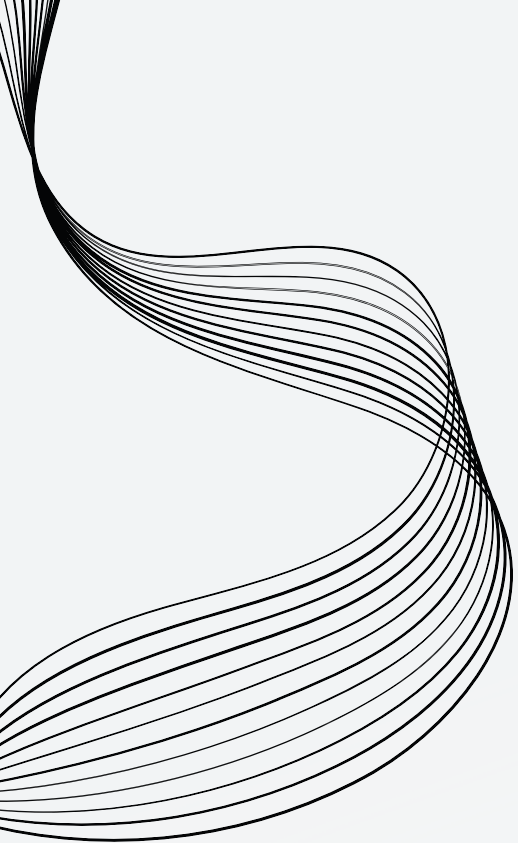
```
[ ] from sklearn.model_selection import train_test_split  
    X_train, X_test, y_train, y_test = train_test_split(X,Y, test_size=0.33, random_state=42)
```

STRATEGY N°3

ANDROID - BASELINE

```
classifiers = {
    "Random Forest": RandomForestClassifier(),
    "Decision Tree": DecisionTreeClassifier(),
    "SVM": SVC(),
    "K-Nearest Neighbors": KNeighborsClassifier(),
    "Logistic Regression": LogisticRegression(),
    "Naive Bayes": GaussianNB(),
    "AdaBoost": AdaBoostClassifier(),
    "Gradient Boosting": GradientBoostingClassifier(),
    "Bagging": BaggingClassifier(),
    "Extra Trees": ExtraTreesClassifier(),
    "Voting": VotingClassifier(estimators=[
        ('lr', LogisticRegression()),
        ('rf', RandomForestClassifier()),
        ('svc', SVC())
    ]),
    "Stacking": StackingClassifier(estimators=[
        ('lr', LogisticRegression()),
        ('rf', RandomForestClassifier()),
        ('svc', SVC())
    ], final_estimator=RandomForestClassifier()),
    "MLP": MLPClassifier(max_iter=1000)
}
```

```
RandomForestClassifier had an accuracy of 0.7452491011813046
DecisionTreeClassifier had an accuracy of 0.6224961479198767
SVC had an accuracy of 0.7170005136106831
KNeighborsClassifier had an accuracy of 0.731381612737545
LogisticRegression had an accuracy of 0.844889573703133
GaussianNB had an accuracy of 0.6640986132511556
AdaBoostClassifier had an accuracy of 0.724191063174114
GradientBoostingClassifier had an accuracy of 0.7652799178222907
BaggingClassifier had an accuracy of 0.7015921931176169
ExtraTreesClassifier had an accuracy of 0.7416538263995891
VotingClassifier had an accuracy of 0.7632254750898819
StackingClassifier had an accuracy of 0.8376990241397021
MLPClassifier had an accuracy of 0.8397534668721109
```



ANDROID - HYPEROPT

```
classifiers = hp.choice(name, [  
    sklearn_DecisionTreeClassifier(),  
    sklearn_KNeighborsClassifier(),  
    sklearn_PassiveAggressiveClassifier(),  
    sklearn_SGDClassifier(),  
    sklearn_XGBClassifier(),  
    sklearn_LinearSVC(),  
    sklearn_SVC(),  
    sklearn_DecisionTreeClassifier(),  
    sklearn_KNeighborsClassifier(),  
    sklearn_PassiveAggressiveClassifier(),  
    sklearn_SGDClassifier(),  
    sklearn_XGBClassifier()  
)  
  
classifier = HyperoptEstimator(classifier=classifiers,  
                               algo=tpe.suggest,  
                               max_evals=40,  
                               trial_timeout=120)  
  
classifier.fit(X_train.values, y_train.values)  
y_pred = classifier.predict(X_test.values)  
accuracy = accuracy_score(y_test.values, y_pred)
```

accuracy: 0.8628659476117103

```
classifier.best_model()  
  
{'learner': SGDClassifier(),  
 'preprocs': (StandardScaler(with_std=False),),  
 'ex_preprocs': ()}
```

EDAIC - BASELINE

train - test - validation → already prepared in datasets

2 baselines → train/test, train/validation

EDAIC - TRAIN/TEST BASELINE

```
classifiers = {  
    "Random Forest": RandomForestClassifier(),  
    "Decision Tree": DecisionTreeClassifier(),  
    "SVM": SVC(),  
    "K-Nearest Neighbors": KNeighborsClassifier(),  
    "Logistic Regression": LogisticRegression(),  
    "Naive Bayes": GaussianNB(),  
    "AdaBoost": AdaBoostClassifier(),  
    "Gradient Boosting": GradientBoostingClassifier(),  
    "Bagging": BaggingClassifier(),  
    "Extra Trees": ExtraTreesClassifier(),  
    "Voting": VotingClassifier(estimators=[  
        ('lr', LogisticRegression()),  
        ('rf', RandomForestClassifier()),  
        ('svc', SVC())  
    ]),  
    "Stacking": StackingClassifier(estimators=[  
        ('lr', LogisticRegression()),  
        ('rf', RandomForestClassifier()),  
        ('svc', SVC())  
    ], final_estimator=RandomForestClassifier()),  
    "MLP": MLPClassifier(max_iter=1000)  
}
```

```
RandomForestClassifier had an accuracy of 0.7747342600163533  
DecisionTreeClassifier had an accuracy of 0.6618969746524939  
SVC had an accuracy of 0.7755519215044971  
KNeighborsClassifier had an accuracy of 0.7334423548650858  
LogisticRegression had an accuracy of 0.7694194603434178  
GaussianNB had an accuracy of 0.42804578904333607  
AdaBoostClassifier had an accuracy of 0.7641046606704824  
GradientBoostingClassifier had an accuracy of 0.7751430907604252  
BaggingClassifier had an accuracy of 0.7661488143908421  
ExtraTreesClassifier had an accuracy of 0.7759607522485691  
VotingClassifier had an accuracy of 0.7755519215044971  
StackingClassifier had an accuracy of 0.7547015535568274  
MLPClassifier had an accuracy of 0.6888798037612428
```

EDAIC - TRAIN/VALIDATION BASELINE

```
RandomForestClassifier had an accuracy of 0.8101774355540676  
DecisionTreeClassifier had an accuracy of 0.6906595246066287  
SVC had an accuracy of 0.8105122196183462  
KNeighborsClassifier had an accuracy of 0.7760294609976565  
LogisticRegression had an accuracy of 0.8028121861399398  
GaussianNB had an accuracy of 0.5038500167392033  
AdaBoostClassifier had an accuracy of 0.7833947104117844  
GradientBoostingClassifier had an accuracy of 0.8068295949112823  
BaggingClassifier had an accuracy of 0.7937730164044191  
ExtraTreesClassifier had an accuracy of 0.8098426514897891  
VotingClassifier had an accuracy of 0.8101774355540676  
StackingClassifier had an accuracy of 0.7820555741546702  
MLPClassifier had an accuracy of 0.792433880147305
```

EDAIC - HYPEROPT

```
classifiers = hp.choice(name, [
    sklearn_DecisionTreeClassifier(),
    sklearn_KNeighborsClassifier(),
    sklearn_PassiveAggressiveClassifier(),
    sklearn_SGDClassifier(),
    sklearn_XGBClassifier(),
    sklearn_LinearSVC(),
    sklearn_SVC(),
    sklearn_DecisionTreeClassifier(),
    sklearn_KNeighborsClassifier(),
    sklearn_PassiveAggressiveClassifier(),
    sklearn_SGDClassifier(),
    sklearn_XGBClassifier()
])

classifier = HyperoptEstimator(classifier=classifiers,
                              algo=tpe.suggest,
                              max_evals=40,
                              trial_timeout=120)

classifier.fit(X_train.values, Y_train.values)
y_pred = classifier.predict(X_vali.values)
accuracy = accuracy_score(Y_vali.values, y_pred)
```

accuracy: **0.7787077335118848**

```
[ ] classifier.best_model()

{'learner': SGDClassifier(),
 'preprocs': (MinMaxScaler(feature_range=(0.0, 1.0)),),
 'ex_preprocs': ()}
```




THANK YOU